1 Introduction

Consider the Billionaire's Problem: Larry and Sergey are both wealthy men. They want to design a protocol to find out whose net worth is higher, without having to reveal their net worths to each other. Since they are business partners, they trust each other to follow the protocol truthfully (i.e., to use their true worth as inputs and to follow the protocol instructions faithfully), but they still do not want the protocol to reveal more about their net worths than is absolutely necessary. Can this be achieved? More fundamentally, how do we define the notion of security for this goal?

Note that "total" privacy cannot be achieved, as they will each learn something new about the other's wealth, namely, whether it is greater or less than his own. In some special cases, Larry can infer Sergey's net worth entirely from just this single piece of knowledge. For example, if Larry's net worth is \$1 and Sergey's net worth is \$0, then when Larry learns that Sergey's worth is less than his own, he can infer that Sergey's net worth is exactly \$0 (we ignore the possibility of Sergey being in debt). However, this leakage of knowledge is inherent in the task they are carrying out for these values, and therefore should not should not be considered a deficiency of any particular protocol they might use.

Alternatively, Larry might start with some prior knowledge about Sergey's wealth that might enable him to infer Sergey's exact net worth from just knowing whose net worth is higher. For example, if Larry knows that Sergey's net worth is either \$4 billion or \$5 billion, and his own worth is \$4.5 billion, then learning who is wealthier immediately reveals Sergey's net worth to Larry. This also should not be considered a violation of our security goal.

In both of these examples, we must be content to let Larry learn whatever he can infer from the final result (i.e., who is wealthier) — but we want him to learn nothing more than that! So in defining security, we will aim to restrict the "relative knowledge" revealed by a protocol, versus what is learned from the outcome alone. Similarly to the setting of zero knowledge, we will do so using the notion of an efficient simulator that is given only the input and output of the party, and must simulate the entire view of that party in the protocol.

2 Secure Two-Party Computation

Here we formalize a model and security definition for the informal goals described above.

2.1 Model

We will consider a very simplified model that does not capture many real-world concerns, but is still rich enough to make the problem interesting and non-trivial.

1. There are two parties (more formally, two ppt algorithms) P_1 and P_2, who have inputs x_1 and x_2 respectively, and wish to evaluate a public polynomial time-computable function $f(\cdot, \cdot)$ on those inputs. For example, in the billionaire's problem, $f(x_1, x_2) = [x_1 > x_2]$.

   Without loss of generality, we may assume that f is a deterministic function that outputs a single value that is given to both parties. If we wish for P_1 and P_2 to receive the outputs of two possibly different deterministic functions $f_1(\cdot, \cdot)$, $f_2(\cdot, \cdot)$ (respectively), this can be emulated using a single function $f'((x_1, r_1), (x_2, r_2)) = (f_1(x_1, x_2) \oplus r_1 \| f_2(x_1, x_2) \oplus r_2)$, where each P_i augments its own input x_i by a uniformly random string r_i of appropriate length. Since r_1

and r_2 are chosen uniformly at random and are independent of everything else, the output $f_1(x_1, x_2)$ is perfectly hidden from P_2, as is $f_2(x_1, x_2)$ from P_1.

We can also evaluate a randomized function f by emulating it with a deterministic function (showing how to do this is one of your homework problems). However, security becomes quite a bit subtler to define in this case; see below.

2. We assume that the parties are semi-honest, often called "honest but curious." That is, they run the protocol exactly as specified (no deviations, malicious or otherwise), but may try to learn as much as possible about the input of the other party from their views of the protocol. Hence, we want the view of each party not to leak more knowledge than necessary.

3. As usual, the view of a party P_i in an interaction with the other party on their inputs x_1, x_2, denoted $\text{view}_{P_i}[P_1(x_1) \leftrightarrow P_2(x_2)]$, consists of its input x_i, the random coins r_{P_i} used by P_i, and all the messages received from the other party. The final output of P_i is denoted $\text{out}_{P_i}[P_1(x_1) \leftrightarrow P_2(x_2)]$.

2.2   Security Definition

Definition 2.1. A pair of ppt machines (P_1, P_2) is a secure 2-party protocol (for static, semi-honest adversaries) for a deterministic polynomial time-computable function $f(\cdot, \cdot)$ if the following properties hold:

1. Completeness: for all $i \in \{1, 2\}$ and all $x_1, x_2 \in \{0, 1\}^*$, we have (with probability 1):

$$\text{out}_{P_i}[P_1(x_1) \leftrightarrow P_2(x_2)] = f(x_1, x_2).$$

2. Privacy: there exist nuppt simulators $\mathcal{S}_1, \mathcal{S}_2$ such that for all $x_1, x_2 \in \{0, 1\}^*$ and all $i \in \{1, 2\}$,

$$\text{view}_{P_i}[P_1(x_1) \leftrightarrow P_2(x_2)] \overset{c}{\approx} \mathcal{S}_i(x_i, f(x_1, x_2)).$$

A few remarks are in order. First, privacy is per-instance: the only knowledge leaked to a party by the protocol on inputs x_1, x_2 is whatever can be inferred (efficiently) from the party's own input and the value of $f(x_1, x_2)$. For example, if the inputs are such that $f(x_1, x_2)$ reveals nothing at all, then the execution of the protocol on those inputs should also reveal nothing; conversely, if the output reveals everything about both parties' inputs, then the protocol is allowed to leak everything as well. Second, any "prior knowledge" that the parties have about each others' inputs is captured by the non-uniformity of the simulator (and implicit distinguisher) in the definition of privacy.

2.3   Definition for Randomized Functions

Definition 2.1 is relatively straightforward due to the simplicities of our model, in particular, the deterministic nature of f. We briefly discuss some of the issues that arise in defining security for randomized functions. First, how should completeness be defined? It no longer makes sense to demand that $\text{out}_{P_1} = f(x_1, x_2)$, since we now have a random variable $f(x_1, x_2; r)$ over the choice of r (which neither party should be able to influence). Instead, we want that both out_{P_1} and out_{P_2} in a single execution of the protocol are simultaneously distributed as $f(x_1, x_2; r)$ for the same random

r. This is so that the protocol between P_1 and P_2 has the effect of emulating a single, consistent randomized evaluation of f. Formally, we want that for all x_1, x_2,

$$(\text{out}_{P_1}, \text{out}_{P_2})[P_1(x_1) \leftrightarrow P_2(x_2)] \stackrel{c}{\approx} (f(x_1, x_2; r), f(x_1, x_2; r)),$$

where r is uniformly random and the same in both appearances of f.

The next natural question is how to define privacy against a semi-honest party. Again, simultaneity of the respective views of P_1 and P_2 is an important issue, and is even more subtle to get right. It turns out that the proper way of addressing all these concerns is to define correctness and privacy all together by comparing two joint distributions: the "real world" distribution of the parties' outputs and the semi-honest party's view, versus the "ideal world" distribution of the function output and simulated view (again, for a single randomized evaluation of f). Formally, we require that there exist nuppt simulators $\mathcal{S}_1, \mathcal{S}_2$ such that for all x_1, x_2 and all $i \in \{1, 2\}$,

$$(\text{out}_{P_1}, \text{out}_{P_2}, \text{view}_{P_i})[P_1(x_1) \leftrightarrow P_2(x_2)] \stackrel{c}{\approx} (f(x_1, x_2; r), f(x_1, x_2; r), S_i(x_i, f(x_1, x_2; r))).$$

Note that the above condition automatically implies the correctness condition above, so it is the only one needed to prove security.

2.4   Secure Protocol for Addition

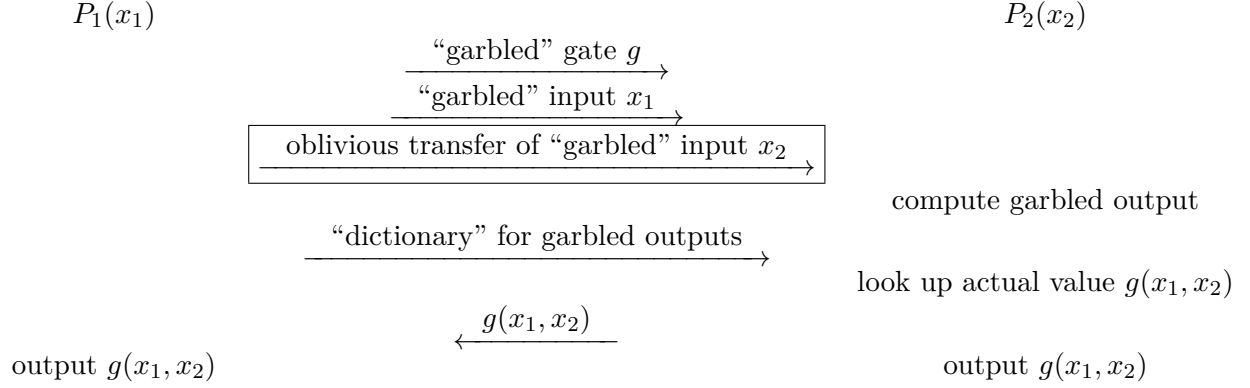As a brief test case, we consider a contrived protocol for evaluating the addition function $f(x_1, x_2) = x_1 + x_2$.

$$P_1(x_1) \qquad\qquad\qquad P_2(x_2)$$
$$\xrightarrow{\quad x_1 \quad}$$
$$\xleftarrow{\quad x_2 \quad}$$
$$\text{output } x_1 + x_2 \qquad\qquad \text{output } x_1 + x_2$$

Clearly the protocol is complete. For privacy, since P_1 is entitled to know the value of $x_1 + x_2$, and also already knows x_1, he can trivially infer the value of x_2. Formally, we can give a simulator $\mathcal{S}_1(x_1, s = f(x_1, x_2) = x_1 + x_2)$ that just outputs the view consisting of input x_1, empty randomness, and a single message $s - x_1$ coming from P_2. Clearly, this view is identical to P_1's view in the real protocol.

3   Secure Protocol for Arbitrary Circuits

We now describe a protocol, originally described by Yao, for evaluating an arbitrary function f represented as a boolean (logical) circuit. We describe the basic idea for just a single logic gate, and then outline how it generalizes to arbitrary circuits.

Let $g : \{0, 1\} \times \{0, 1\} \to \{0, 1\}$ be an arbitrary logic gate on two input bits (e.g., the NAND function). Party P_1 holds the first input bit $x_1 \in \{0, 1\}$ and party P_2 holds the second input bit $x_2 \in \{0, 1\}$. Together they wish to compute $g(x_1, x_2)$ securely, in the sense of Definition 2.1.

At a high level, the protocol works like this:

$P_1(x_1)$ $\qquad\qquad\qquad\qquad\qquad\qquad\qquad\qquad\qquad\qquad\qquad$ $P_2(x_2)$

$$\xrightarrow{\text{“garbled” gate } g}$$

$$\xrightarrow{\text{“garbled” input } x_1}$$

$$\xrightarrow{\text{oblivious transfer of “garbled” input } x_2}$$

compute garbled output

$$\xrightarrow{\text{“dictionary” for garbled outputs}}$$

look up actual value $g(x_1, x_2)$

$$\xleftarrow{g(x_1, x_2)}$$

output $g(x_1, x_2)$ $\qquad\qquad\qquad\qquad\qquad\qquad\qquad\qquad\qquad$ output $g(x_1, x_2)$

For intuition, the crucial points for security are:

- P_2 never sees x_1 in an "ungarbled" form, so P_2 learning nothing about x_1.

- By the security of the "oblivious transfer" sub-protocol (described below), P_1 learns nothing about x_2.

- Using the garbled inputs x_1, x_2 with the garbled gate, P_2 can "obliviously" compute the garbled output for only the correct value of $g(x_1, x_2)$.

Concretely, these ideas are implemented using basic symmetric-key encryption. The idea is the following: each "wire" of the gate (the two inputs and one output) is associated with a pair of random symmetric encryption keys, chosen by P_1; the two keys correspond to the two possible "values" (0 or 1) that the wire can take. For $i = 1, 2$, let k_0^i, k_1^i be the keys corresponding to the input wire x_i, and let k_0^o, k_1^o be the keys corresponding to the output wire. The "garbled circuit" that P_1 sends to P_2 is a table of four doubly encrypted values, presented in a random order:

$$
\boxed{
\begin{array}{l}
\mathsf{Enc}_{k_0^1}(\mathsf{Enc}_{k_0^2}(k_{g(0,0)}^o)) \\
\mathsf{Enc}_{k_0^1}(\mathsf{Enc}_{k_1^2}(k_{g(0,1)}^o)) \\
\mathsf{Enc}_{k_1^1}(\mathsf{Enc}_{k_0^2}(k_{g(1,0)}^o)) \\
\mathsf{Enc}_{k_1^1}(\mathsf{Enc}_{k_1^2}(k_{g(1,1)}^o))
\end{array}
}
$$

Observe that if P_2 knows (say) k_0^1 and k_1^2, i.e., the keys corresponding to inputs $x_1 = 0$ and $x_2 = 1$, then P_1 can decrypt $k_{g(0,1)}^o$, the key corresponding to the output value of the gate, but none of the other entries! (Note that this requires the encryption scheme to satisfy some simple properties, such as the ability to detect when a ciphertext has decrypted successfully. These properties are easy to obtain.) The random order of the table prevents P_1 from learning the "meaning" of the keys that it knows, otherwise this information would be leaked by which of the table entries decrypt properly. In conclusion, knowing exactly one key for each input wire allows P_2 to learn exactly one key (the correct one) for the output wire, without learning the meanings of any of the keys.

The only remaining question is how P_2 obtains the right keys for the input wires. For x_1, this is easy: P_1 just sends P_2 the key $k_{x_1}^1$ corresponding to its input bit x_1. Note that P_2 learns nothing about x_1 from this. Next, P_1 and P_2 run an "oblivious transfer" protocol (described in the next subsection) which allows P_2 to learn $k_{x_2}^2$, and only $k_{x_2}^2$, without revealing anything about the value of x_2 to P_1.

4

Finally, P_1 tells P_2 the "meanings" of the two possible output keys k_0^o, k_1^o, which reveals to P_2 the value of $g(x_1, x_2)$. Then P_2 sends this value to P_1 as well (recall that both parties are semi-honest, so neither will lie).

For more complex circuits f, the protocol generalizes in a straightforward manner: P_1 chooses two keys for every wire in the circuit, and constructs a garbled table for each gate, using the appropriate keys for the input and output wires. P_2 can compute the garbled gates iteratively, while remaining oblivious to the meanings of the intermediate wires. Then P_1 finally reveals the meanings of just the output wires.

A proof of security for this construction is beyond the scope of this lecture, but contains no surprises; see the paper by Lindell and Pinkas for a full rigorous proof. The key point is that a simulator can construct garbled gates that always result in the same key being decrypted (irrespective of which inputs keys were used), thus allowing the simulator to "force" P_2 to output the correct value $f(x_1, x_2)$. Security of the symmetric encryption scheme prevents P_2 from detecting these malformed garbled gates, since P_2 can only decrypt one entry from each gate.

4 Oblivious Transfer

We conclude by describing how to perform an oblivious transfer between the two parties. We will consider the specific form of oblivious transfer that is required to complete our protocol: P_1 is holding two bits b_0, b_1 and wants to transfer exactly one of them to P_2, according to P_2's choice bit $\sigma = x_2$, while learning nothing about which one was received. (To transfer entire keys $k_0, k_1 \in \{0,1\}^n$, the parties can just run the protocol n times, using the same choice bit σ and the corresponding pairs of bits from k_0, k_1).

Our protocol relies on a family of trapdoor permutations $\{f_s \colon \{0,1\}^n \to \{0,1\}^n\}$ with hard-core predicate $h \colon \{0,1\}^n \to \{0,1\}$.

$$
\begin{array}{ll}
P_1(b_0, b_1) & P_2(\sigma) \\
\text{choose } (f_s, f_s^{-1}) & \\
& \xrightarrow{\quad f_s \quad} \\
& v_\sigma \leftarrow \{0,1\}^n \\
& w_\sigma = f_s(v_\sigma) \\
& w_{1-\sigma} \leftarrow \{0,1\}^n \\
& \xleftarrow{\quad w_0, w_1 \quad} \\
\text{let } v_0 = f_s^{-1}(w_0),\ v_1 = f_s^{-1}(w_1) & \\
\xrightarrow{\quad c_0 = h(v_0) \oplus b_0,\ c_1 = h(v_1) \oplus b_1 \quad} & \\
& \text{output } h(v_\sigma) \oplus c_\sigma
\end{array}
$$

In words, P_1 picks a random function f_s (with trapdoor) from the family and sends it to P_2. Then P_2 chooses uniformly random $w_0, w_1 \in \{0,1\}^n$ so that it knows the preimage of only w_σ, and sends these to P_1. (Observe that this reveals no information about σ to P_1 because w_0, w_1 are uniform and independent.) Next, P_1 encrypts its two bits b_0, b_1 using the hard-core predicate h on the preimages of w_0, w_1, respectively. Finally P_2, knowing the preimage w_σ, can decrypt b_σ, but it learns nothing about $b_{1-\sigma}$ due to the hardness of h. Note that the protocol crucially relies on the fact that P_2 is semi-honest, otherwise it could choose w_0, w_1 so that it knew both preimages. (A full, formal proof of security for this protocol is not too hard, and is a worthwhile exercise.)